

1 **Cross-Cohort Generalization Failure in Gut Microbiome-Based Alzheimer’s**
2 **Disease Classifiers: Evidence from Five Independent Cohorts with Batch**
3 **Correction Analysis**

4 **Running title:** Microbiome AD classifier generalization failure

5 **Authors:** Arhan Chakravarthy^{1*}

6 **Affiliations:** ¹ Independent Researcher, Cupertino, CA, United States

7 **Correspondence:** arhanc360@gmail.com

8 **Keywords:** gut microbiome; Alzheimer’s disease; machine learning; cross-cohort generalization; batch
9 correction; SHAP; PERMANOVA; 16S rRNA sequencing

10 **Figures:** 6 · **Tables:** 4

11 **Abstract**

12 Gut microbiome dysbiosis has been proposed as an Alzheimer’s disease (AD) biomarker, but the
13 transportability of microbiome classifiers across independent cohorts remains uncertain. We reprocessed
14 five public 16S rRNA cohorts (509 samples; China, Kazakhstan, South Korea) with a common
15 DADA2/genus-level pipeline and evaluated logistic regression and LightGBM classifiers using nested
16 within-cohort cross-validation, leave-one-cohort-out (LOCO) transfer, pairwise transfer, PERMANOVA,
17 batch adjustment, and out-of-fold SHAP analyses. Within four labeled cohorts, within-cohort AUC-ROC
18 ranged from 0.633 to 0.998. The near-perfect Zhu 2022 result was stable across 50 repeated nested-CV
19 partitions (mean AUC 0.9955 for logistic regression and 0.9760 for LightGBM) and exceeded a 1,000-
20 permutation label null for both models (empirical p=0.000999). In LOCO evaluation, AUC ranged from
21 0.504 to 0.813, with mean decreases of 0.176 for logistic regression and 0.146 for LightGBM; paired
22 diagnosis-stratified bootstrap intervals excluded zero for five of eight cohort-model differences. Re-
23 deriving genus selection strictly from training data changed LOCO AUC by at most 0.066. In the four-
24 cohort marginal PERMANOVA model, cohort-associated compositional variation exceeded diagnosis-
25 associated variation ($R^2=0.172$ versus 0.014), but this contrast decreased to 0.059 versus 0.023 when
26 restricted to three more technically comparable Chinese paired-end MiSeq cohorts, and dispersion differed
27 significantly among cohorts. Label-blind transductive ComBat-seq and MMUPHin adjustment changed
28 mean LOCO AUC by at most 0.007 and did not recover transferability. Cross-cohort overlap of top
29 predictive taxa was low: logistic-regression overlap remained above a prevalence-restricted random-set
30 null, whereas LightGBM overlap did not. A descriptive coefficient-stability screen identified eight logistic-
31 regression genera with opposite stable fitted-coefficient directions across cohorts. These findings
32 demonstrate substantial, heterogeneous cross-cohort generalization limitations among the evaluated public
33 16S datasets and show that the two tested batch-adjustment methods did not restore transferability.
34 Prospective validation in larger, geographically and technically diverse cohorts is needed before
35 microbiome-based AD classifiers can be considered clinically actionable.

1. Introduction

Alzheimer’s disease (AD) is the most common cause of dementia, accounting for an estimated 60–70% of cases; approximately 57 million people worldwide were living with dementia in 2021 (World Health Organization, 2026). The gut–brain axis has emerged as a mechanistic pathway linking intestinal microbial communities to neurodegeneration: altered gut microbiome composition has been associated with AD pathology through neuroinflammatory, metabolic, and neuroendocrine routes (Cryan et al., 2019; Kowalski and Mulak, 2019). Since the landmark observation that germ-free APPPS1 transgenic mice showed reduced amyloid- β pathology in the absence of gut microbiota compared with conventionally colonized controls (Harach et al., 2017), dozens of 16S rRNA and shotgun metagenomic studies have reported significant differences in gut microbial composition between AD or MCI patients and cognitively normal (CN) controls. MCI is included as a diagnostic category because it can represent a preclinical or prodromal stage of AD, though not all MCI progresses to AD, and the primary supervised analyses here focus on binary AD versus CN comparisons (Section 2.4).

The clinical appeal of a gut microbiome-based diagnostic signature is considerable: stool is non-invasive, inexpensive to collect, and microbiome profiling is increasingly accessible. Multiple independent groups have trained machine learning classifiers on gut microbiome abundance profiles and reported promising within-cohort discrimination, with AUC-ROC of 0.70–0.90 in several published studies (Zhuang et al., 2018; Ling et al., 2021; Zhu et al., 2022). A critical property of any diagnostic biomarker — transferability across independent populations — has rarely been tested in this literature.

The limited cross-cohort comparisons that do exist reveal a striking inconsistency in the direction of reported associations. A systematic review and meta-analysis synthesizing 17 published studies (679 AD/MCI patients, 632 controls) documented that *Bacteroides* abundance was significantly elevated in North American cohorts yet significantly reduced in Chinese cohorts — a directional reversal across geographically distinct populations (Jemimah et al., 2023). Similar contradictions have been observed for *Akkermansia*, *Faecalibacterium*, and *Lactobacillus* (Vogt et al., 2017; Liu et al., 2019; Zhuang et al., 2018), suggesting published within-cohort classifiers may detect cohort-specific compositional signatures — driven by diet, geography, host genetics, sequencing protocols, or a combination — rather than a universal microbial correlate of AD pathology.

A recent multi-cohort study integrating 15 public AD microbiome projects identified candidate gut microbiota indicators across multiple geographic populations (Wang et al., 2026), confirming multi-population integration remains a methodological frontier. The present study is complementary, systematically quantifying cross-cohort generalization failure and empirically evaluating whether two commonly used batch-correction approaches recover cross-cohort transferability — questions that remain open even where multi-cohort classifiers have been attempted.

Despite this recognized heterogeneity, few studies have jointly quantified: (1) the magnitude of cross-cohort AUC collapse when trained classifiers are applied to independent cohorts; (2) the relative magnitude of cohort- versus diagnosis-associated compositional variance, as estimated by variance partitioning — noting “cohort” is itself a composite, study-of-origin variable (bundling sequencing platform, library preparation, recruitment site, diagnostic ascertainment, geography, and diet, among other differences) that a method such as PERMANOVA cannot decompose into technical and biological components; and (3) whether the two evaluated batch-adjustment methods can recover cross-cohort

generalization. Answering these three questions jointly in one rigorous analysis is the primary contribution of this paper.

We assembled five publicly available 16S rRNA gene sequencing cohorts spanning three countries (China, Kazakhstan, South Korea) and conducted a comprehensive cross-cohort generalization evaluation. We report that classifiers showed substantial and variable degradation when transferred to independent cohorts; cohort identity explains substantially more compositional variance than disease status, though the size of this disparity depends materially on which cohorts are compared; under label-blind transductive adjustment, ComBat-seq and MMUPHin left mean cross-cohort AUC essentially unchanged and did not recover transferability; and top predictive taxa showed low absolute overlap across cohorts, with a subset showing opposite, stable fitted-coefficient directions. Together, these findings constitute a systematic quantification of cross-cohort generalization limitations among the evaluated public gut microbiome AD classifier datasets, with direct implications for the field’s trajectory toward clinical translation.

2. Materials and Methods

2.1 Cohort Identification and Data Acquisition

We performed a systematic search of NCBI SRA and ENA for publicly available 16S rRNA amplicon sequencing studies of gut microbiota in AD, MCI, or cognitively normal populations, deposited before January 2025, using the following SRA search terms: (“Alzheimer’s disease” OR “Alzheimer disease” OR “mild cognitive impairment” OR “MCI” OR “dementia”) AND (“gut microbiome” OR “gut microbiota” OR “intestinal microbiome” OR “16S rRNA”) AND (“Homo sapiens”). Studies were included if they provided (1) per-sample diagnosis labels (AD, MCI, or CN) in public metadata, (2) raw FASTQ files or quality-filtered reads in a public archive, and (3) minimum sample sizes of 40 participants per cohort. We restricted this analysis to 16S rRNA studies; shotgun metagenomic cohorts were identified but excluded to avoid the additional platform confound of mixing amplicon and whole-genome sequencing data. Cohorts were excluded if per-sample labels were IRB-restricted (for supervised analyses), data had not been publicly released at the time of analysis, or samples represented a tissue compartment other than fecal microbiome. A complete record of the search procedure, identified cohorts, and exclusion rationale is in Supplementary Note S1.

Four cohorts met all criteria for supervised analysis. A fifth cohort, Kim/KBASE 2022 (PRJEB50447), lacked publicly available per-sample diagnosis labels and was retained only for unsupervised cohort-level analyses. The ALBION 2025 Greek cohort (PRJNA1297934) was identified but excluded because data had not yet been publicly released at the time of analysis. Group-level counts from the published paper (18 amyloid-positive preclinical AD; 60 cognitively normal) were available but insufficient for supervised classifier training or evaluation.

The four cohorts with per-sample diagnosis labels used in supervised analyses were:

Zhuang 2018 (China; PRJNA554111): 86 participants (43 AD, 43 CN) from the Third Military Medical University, Chongqing. Paired-end MiSeq 2×300 bp, V3–V4 region (Zhuang et al., 2018).

Ling 2020 (China; PRJNA633959): 171 participants (100 AD, 71 CN). Paired-end MiSeq 2×300 bp, V3–V4 region (Ling et al., 2021).

Zhu 2022 (China; PRJNA489760): 180 samples deposited in SRA from a larger 302-participant study, representing 90 participants (30 AD, 30 MCI, 30 CN), each contributing two samples: a fecal (*F_*) sample and a blood microbiome (*B_*) sample. Only the 90 fecal samples were used in all analyses; blood samples were excluded as a distinct biological compartment that would confound within-cohort classifiers. Binary analyses used 30 AD and 30 CN fecal samples (n=60). Paired-end MiSeq 2×300 bp, V3–V4 (Zhu et al., 2022). Diagnostic ascertainment differed across cohorts: Zhuang 2018, Ling 2020, and Zhu 2022 used clinical AD diagnosis; Kazakhstan used NINCDS-ADRDA criteria (confirmed by neurologists); and Kim/KBASE 2022 used amyloid-PET positivity for preclinical AD. These differences in diagnostic ascertainment are acknowledged as a source of between-cohort biological heterogeneity.

Kazakhstan 2022 (PRJNA811324): 84 participants (41 AD, 43 CN). Single-end NovaSeq sequencing, V3–V4 region (Kaiyrylkyzy et al., 2022) — the only single-end dataset in the study; implications for read length and taxonomic resolution are discussed in Section 2.2 and Limitations.

The Kim/KBASE 2022 cohort (Jung et al., 2022; PRJEB50447; 78 participants; 18 preclinical AD by amyloid-PET, 60 CN) was included in PERMANOVA analyses (Phase 4) only and excluded from all supervised classifier training and evaluation.

All raw FASTQ data were downloaded from NCBI SRA using fasterq-dump (SRA Toolkit v3.0.10). Cohort-level sequencing statistics are reported in Supplementary Table S2.

2.2 DADA2 Amplicon Sequence Variant Processing

Raw paired-end reads (four cohorts) and single-end reads (Kazakhstan) were processed independently per cohort using DADA2 v1.28 in R v4.3 (Callahan et al., 2016), to avoid introducing cross-cohort data leakage during error model estimation.

For paired-end cohorts, reads were trimmed to truncation lengths selected per cohort based on per-base quality score profiles (generally 240 bp forward / 200 bp reverse for V3–V4 MiSeq 2×300 data), filtered to remove reads with >2 expected errors, chimera-removed using the consensus method, and merged. For Kazakhstan (single-end), the single-end DADA2 workflow was applied without merging; truncation length was adjusted to 250 bp after inspection of quality profiles (Supplementary Note S2). The mixed single-end/paired-end design is an additional technical source of between-cohort heterogeneity acknowledged as a study limitation.

Taxonomic assignment used the SILVA reference database v138.1. ASVs were collapsed to genus level by summing raw counts per sample. A genus was retained across the study if it reached $\geq 20\%$ prevalence (nonzero count in $\geq 20\%$ of samples) in at least one of the five cohorts; this threshold was computed per cohort and the union of genera passing in any cohort was taken across all five, producing a unified feature set of 396 genera. Genera absent in a given sample received a count of zero. Because this feature-selection step used all five cohorts’ raw prevalence simultaneously, the resulting 396-genus feature universe was informed, in part, by the prevalence of the eventual test cohort in every downstream leave-one-cohort-out (LOCO) and pairwise-transfer experiment (Section 2.5). This is not a leakage-free design; it is identified transparently as a global, all-cohort feature-selection step, and its effect on the paper’s conclusions is evaluated directly by a strict training-only sensitivity analysis (Section 2.5.1) rather than assumed negligible.

Leakage status of every preprocessing step, stated explicitly. (1) *DADA2 ASV inference and error-model estimation* (this section) were performed once independently within each cohort using that cohort’s

full read set, with no diagnosis labels and no samples from other cohorts, and were not re-estimated within outer-CV folds. Thus, this raw-read stage does not introduce diagnosis-label leakage or cross-cohort pooling, although within-cohort performance is not fully fold-isolated at the denoising stage because error-model estimation used the cohort's full read set. The strict training-only sensitivity below re-derives genus selection and downstream supervised preprocessing/model fitting, not DADA2. (2) *Global genus-level feature selection* ($\geq 20\%$ prevalence, this section) pools raw prevalence across all five cohorts, including the eventual held-out cohort in every LOCO/pairwise experiment — a genuine, acknowledged limitation bounded by the training-only sensitivity in Section 2.5.1, which re-derives the feature set from training-side samples only and finds the qualitative conclusion unchanged. (3) *Zero-replacement and CLR transformation* (Section 2.3) are deterministic functions of each sample's own count vector, using no information from any other sample, cohort, or diagnosis label. (4) *Model tuning and fitting* (Sections 2.4–2.5) use only the training-side samples of each split, with test-side samples never used for tuning, fitting, or standardization. (5) The *strict training-only sensitivity* (Section 2.5.1) additionally re-derives genus selection from training-side samples only, closing the pathway identified in (2) for that analysis.

2.3 Compositional Data Transformation

Raw genus-level count tables were converted to centered log-ratio (CLR) coordinates for all downstream analyses requiring distributional assumptions on unconstrained data (machine learning classifiers, PERMANOVA on Euclidean distance). The CLR transformation is defined as:

$$\text{clr}(x_i) = \log\left(\frac{x_i}{g(x)}\right)$$

where $g(x)$ is the geometric mean of the composition x .

Zero counts were replaced prior to transformation using the multiplicative replacement strategy (Martín-Fernández et al., 2003) with $\delta = 0.65$:

$$x_i^* = \begin{cases} \delta & \text{if } x_i = 0 \\ (1 - n_0 \cdot \delta / T) \cdot x_i & \text{if } x_i > 0 \end{cases}$$

where $T = \sum_j x_j$ is the total read count of the sample and n_0 is the count of zeros in the sample. The unified genus matrix contained 75.4% zero entries (151,926 of 201,564 cells across the 509 fecal samples and 396 genera). The value $\delta = 0.65$ was chosen such that the replacement mass remained small relative to typical per-sample sequencing depth (median 29,521 reads): in the worst-observed sample, the total mass allocated to zero-filled entries was 2.53% of that sample's total reads. A built-in safeguard rescales δ downward whenever $n_0 \cdot \delta \geq T$, capping replacement mass at 1% of sample total in such cases. This strategy preserves the ratios among non-zero components while introducing a minimal perturbation proportional to the relative frequency of zeros, and is preferable to the commonly used additive pseudocount which distorts low-abundance taxa more severely.

CLR-transformed values were used as features for all machine learning models and for Aitchison distance computation (Euclidean distance on CLR values). Raw genus-level count matrices (not CLR-transformed) were provided as input to ComBat-seq, which requires non-negative integer counts. Relative abundances (proportions) were provided to MMUPHin. Bray-Curtis dissimilarity was computed on relative abundances. CLR values were used for all other analyses.

2.4 Within-Cohort Classifier Baselines

We trained binary classifiers (AD vs. CN) independently for each of the four labeled cohorts using two model families:

L2-regularized logistic regression (scikit-learn v1.5; Pedregosa et al., 2011; LogisticRegression, solver="saga"). Regularization strength C was tuned over {0.001, 0.01, 0.1, 1.0, 10.0} in inner cross-validation. An `l1_ratio` parameter is present in the estimator configuration but has no effect under the L2 penalty actually used (scikit-learn's `penalty` argument was not set to "elasticnet"); all logistic-regression results in this paper are L2 (ridge)-regularized, not elastic-net.

LightGBM gradient-boosted trees (lightgbm v4; Ke et al., 2017). Hyperparameters tuned: number of estimators $\in \{50, 100, 200\}$, learning rate $\in \{0.05, 0.1\}$, max depth $\in \{3, 5\}$, number of leaves $\in \{15, 31\}$.

For Zhu 2022, AD was defined as the positive class and MCI samples were excluded from binary analyses (resulting in $n=60$: 30 AD, 30 CN fecal samples); a separate three-class analysis (AD / MCI / CN) used all 90 fecal samples and is reported as a secondary finding. Because only fecal samples were retained, each patient appears exactly once; standard StratifiedKFold was used without any paired-sample grouping structure.

Performance was estimated using nested cross-validation: a 10-fold outer loop for performance estimation and a 5-fold inner loop for hyperparameter tuning. The primary metric was AUC-ROC computed on out-of-fold predictions pooled across the 10 outer folds, providing a single unbiased AUC estimate per cohort-model combination. Within-cohort 95% CIs (Table 2, Supplementary Figure S1) use the diagnosis-stratified, participant-level percentile bootstrap described in Section 2.5.2 (10,000 replicates); other CIs in this paper use the non-stratified 1,000-resample bootstrap described in Section 2.5, as noted at each occurrence. We also report sensitivity and specificity at the Youden's J optimal threshold, and note published AUC values from the original cohort papers in Table 2 where available.

2.4.1 Zhu 2022 Robustness: Repeated Cross-Validation and Label Permutation

Because Zhu 2022's near-perfect within-cohort AUC (Section 3.2) is estimated from only 60 samples (30 AD, 30 CN), we assessed its robustness with two additional analyses under the strict training-only feature-selection pipeline (Section 2.5.1), so that genus selection, model tuning, and model fitting used no held-out-fold information (DADA2 denoising, Section 2.2, is a per-cohort, label-free step performed once and is not re-derived per fold).

Repeated nested cross-validation. The 10-fold outer / 5-fold inner nested-CV procedure (Section 2.4) was repeated 50 times under different pre-specified deterministic seeds (1000–1049) for the stratified outer-fold partition, with genus selection, zero-replacement, CLR transformation, and model tuning re-derived independently within every outer fold using only that fold's training-side samples. We report the mean, SD, median, 2.5th/97.5th percentile, minimum, and maximum pooled-OOF AUC across the 50 repetitions. This quantifies sensitivity to fold assignment, not independent replication — the same 60 participants underlie all 50 repetitions.

Label permutation test. Under the same pipeline and fixed outer-fold partition (`random_state=42`): for each of exactly 1,000 permutations, AD/CN labels were shuffled across the 60 participants while preserving the 30/30 class balance. Within every outer fold, all label-dependent model fitting and hyperparameter tuning were rerun using the permuted training labels. Fold-specific genus selection, zero

replacement, and CLR construction depend only on fold membership and feature counts rather than diagnosis labels; these components were therefore precomputed once for the fixed partition and reused across permutations. This is an implementation optimization and does not alter the permutation null because the classifier is refit on the permuted labels in every replicate. Empirical $p = (1 + \#\{\text{null AUC} \geq \text{observed AUC}\}) / 1001$, using the training-only single-partition AUC as the observed value; the minimum attainable p at this permutation count is $1 / 1001 \approx 0.000999$, and we do not interpret any p -value below this floor as more precise than it.

2.5 Cross-Cohort Generalization Experiment

We evaluated two cross-cohort transfer designs:

Leave-one-cohort-out (LOCO): For each test cohort, a classifier was trained on the pooled samples from the remaining three labeled cohorts (no cross-validation within this training step; hyperparameters were selected via 5-fold inner CV on the training pool), then evaluated on the held-out cohort, producing four LOCO AUC estimates per model type. We report the LOCO AUC drop as the difference between within-cohort (nested CV) AUC and LOCO AUC. Uncorrected LOCO 95% CIs (Table 3, Figure 3A) use the diagnosis-stratified, participant-level bootstrap of Section 2.5.2 (10,000 replicates); CIs on batch-corrected LOCO AUC (Table 4, Figure 5, Supplementary Figure S4) use the non-stratified 1,000-resample bootstrap described below. Non-overlap of two independently computed CIs is not a formal statistical test of whether their difference is nonzero and is not used as one anywhere in this paper; the AUC drop is instead evaluated with the paired bootstrap procedure in Section 2.5.2.

Pairwise transfer: A classifier was trained on one cohort and tested on each of the other three individually, yielding a 4×4 matrix (diagonal excluded) of cross-cohort AUC per model type — more stringent than LOCO and able to isolate specific directional failures. Pairwise AUC 95% CIs (Supplementary Table S1) use the non-stratified 1,000-resample bootstrap on test-set predictions, not the diagnosis-stratified regeneration of Section 2.5.2, which covers only the within-cohort and uncorrected-LOCO CIs.

Sensitivity analyses: To assess whether generalization failure was attributable to any single cohort, we repeated LOCO under three additional configurations: (A) excluding Kazakhstan entirely, LOCO evaluated on the remaining three; (B) excluding Zhuang 2018; (C) excluding Ling 2020. Results are in Section 3.3 and Supplementary Figure S3.

2.5.1 Training-Only Feature-Selection Sensitivity Analysis

Because the 396-genus feature set (Section 2.2) was selected globally using all five cohorts’ raw prevalence, every LOCO and pairwise experiment above used a feature universe partly informed by the eventual held-out cohort. To test whether this affected the generalization-failure conclusion, we repeated the within-cohort, LOCO, and pairwise-transfer analyses with genus selection re-derived independently inside every split, using only that split’s training-side samples: for each outer fold (within-cohort), held-out cohort (LOCO), and train cohort (pairwise), the $\geq 20\%$ -prevalence filter was recomputed from training-side AD/CN-labeled samples only, and test/held-out samples were projected onto the training-derived genus set (absent genera zero-filled) before per-sample zero-replacement and CLR transformation. Model families, hyperparameter grids, random state, fold counts, and bootstrap CI procedure matched the main analyses (Sections 2.4–2.5). This yields a strictly train-only feature universe per split, at the cost of a smaller, split-

specific genus count (e.g., 150 genera when Kazakhstan is excluded during its own LOCO hold-out, versus 396 globally).

2.5.2 Formal Within-Cohort vs. LOCO AUC Difference

Non-overlapping bootstrap CIs are a heuristic, not a formal statistical comparison. To formally evaluate the within-cohort-versus-LOCO AUC difference for each of the eight cohort–model combinations, we performed a paired, diagnosis-stratified, participant-level bootstrap: for each pair, the within-cohort and LOCO predictions refer to the same participants, and each of 10,000 replicates resampled participants with replacement (separately within AD and CN) and used that same resampled index set for both AUC values — so the resulting distribution of $\Delta = \text{AUC}_{\text{within}} - \text{AUC}_{\text{LOCO}}$ is a genuine paired comparison, not the difference of two independently resampled statistics. We report the observed delta, bootstrap mean delta, 95% percentile CI, and the proportion of replicates with $\Delta \leq 0$ (an empirical bootstrap quantity, not an exact p-value). Within-cohort and uncorrected-LOCO CIs in Tables 2 and 3 (Figures 2 and 3A, Supplementary Figure S1) were regenerated with this same procedure, so every within-cohort and uncorrected-LOCO CI in the paper uses one consistent resampling method; this did not extend to the pairwise-transfer CIs (Supplementary Table S1) or batch-corrected LOCO CIs (Table 4, Figure 5, Supplementary Figure S4), which retain the original non-stratified 1,000-resample bootstrap of Section 2.5 — each table/figure states which procedure its CIs use.

Within-cohort out-of-fold prediction files lack an explicit participant identifier; identity was recovered by replicating the exact data-loading/filtering sequence and verified, before use, by an exact-match assertion against the file’s own saved outcome column, for every cohort–model combination.

2.6 PERMANOVA Variance Decomposition

To quantify the relative contribution of cohort identity and disease status to total compositional variance, we conducted PERMANOVA using the `adonis2` function in the R package `vegan` v2.6 (Oksanen et al., 2024), with two distance matrices: **Aitchison distance** (Euclidean distance on CLR values), the compositionally appropriate distance that does not distort due to the simplex constraint, and **Bray-Curtis dissimilarity** (on relative abundances), included for comparability with the broader microbiome literature.

All PERMANOVA models used 9,999 permutations with Type III (marginal) sums of squares (by=“margin”) for independent assessment of cohort and diagnosis effects; the minimum attainable permutation p-value is reported as $p=0.0001$ throughout. Three models were fitted: (1) cohort only, all five cohorts ($n=509$); (2) cohort + diagnosis, restricted to the four labeled cohorts ($n=401$: AD and CN samples only, Zhu 2022 MCI excluded), with marginal R^2 per term; (3) within each labeled cohort, diagnosis as sole predictor (within-cohort disease effect size). Figure 4B displays the PCoA from Model 1 ($n=509$, including MCI samples), while marginal R^2 values reported throughout the text derive from Model 2 ($n=401$).

Homogeneity of multivariate dispersions was tested using `betadisper` and `permutest` (9,999 permutations) to assess whether observed cohort differences reflected centroid separation, within-group spread heterogeneity, or both; a significant beta-dispersion test indicates PERMANOVA R^2 partially reflects within-group spread, not only centroid shifts. To identify which cohort pairs drive the overall dispersion difference, we additionally computed all pairwise `betadisper` comparisons among the five cohorts

(Aitchison distance, 9,999 permutations per pair), with Holm and Benjamini-Hochberg multiplicity adjustment across the 10 pairwise p-values.

Cohort is a composite variable; PERMANOVA cannot decompose it into technical and biological sources. “Cohort” in this study bundles sequencing platform, library preparation, recruitment site and country, clinical diagnostic ascertainment, and unmeasured dietary and lifestyle differences into a single categorical variable. A PERMANOVA or betadisper term for cohort therefore estimates the total compositional variance associated with study-of-origin, not a technical-batch-effect estimate specifically; no method here decomposes that total into technical and biological components. To assess how sensitive the cohort-versus-diagnosis variance disparity (Section 3.4) is to this composite structure, we repeated the marginal PERMANOVA and betadisper analyses on a subset restricted to three more technically comparable cohorts — Zhuang 2018, Ling 2020, and Zhu 2022, all Chinese, paired-end MiSeq 2×300bp, V3–V4 studies — while still differing in recruitment site, exact diagnostic procedure, and other unmeasured factors, so this subset is more comparable, not identical. This sensitivity used the same marginal (`by="margin"`) two-variable PERMANOVA specification and 9,999 permutations as the main four-cohort model, restricted to binary AD/CN samples from the three cohorts ($n=317$), with betadisper/permutest and pairwise post-hoc comparisons on the same subset.

To further quantify how much cohort-specific information is retained in the preprocessed feature space, we trained a logistic regression classifier (one-vs-rest, stratified 5-fold CV) to predict cohort identity from the CLR-transformed features across all 509 samples; the macro-AUC provides a complementary measure of cohort separability.

2.7 Batch Correction

We evaluated one count-based batch-correction method (ComBat-seq) and one microbiome-specific meta-analysis framework (MMUPHin), each with cohort identity as the batch variable, under two designs differing in whether diagnosis information is supplied during correction. Both designs are **transductive**: both methods estimate batch-effect parameters jointly across every sample being corrected — including, in the downstream LOCO evaluation, the eventual held-out test cohort’s own feature values and cohort membership. Neither is prospective or inductive (neither defines a transformation applicable to a genuinely new, unseen patient without refitting); they differ only in whether the target cohort’s diagnosis labels are additionally used.

Design 1 — label-blind transductive correction (primary). ComBat-seq (Zhang et al., 2020) was run on raw genus-level count matrices ($\text{genera} \times \text{samples}$, four labeled cohorts) with `group = NULL`, `full_mod = FALSE` — no diagnosis information supplied. MMUPHin v1.15 (Ma et al., 2022) was run on relative-abundance profiles with `batch = cohort`, `covariates = NULL`; the metadata object given to `adjust_batch` included no diagnosis column. Both use cohort identity and every sample’s own raw feature values (making the design transductive) but never a true AD/CN label. Corrected tables were CLR-transformed using the zero-replacement procedure of Section 2.3.

Design 2 — label-informed transductive sensitivity analysis (exploratory). The same methods were additionally run with `group = disease_status`, `full_mod = TRUE` (ComBat-seq) and `covariates = diagnosis` (MMUPHin), using every sample’s own true diagnosis label — including the eventual held-out-cohort’s labels in the LOCO evaluation. Because correction is fit using test-cohort samples’ own diagnosis labels, this design cannot estimate prospective external-validation performance

and is retained only as an exploratory sensitivity analysis, not as a performance bound: access to test labels during correction can improve or worsen downstream classifier performance, so it is not mathematically guaranteed to bound what label-blind correction could achieve.

After each correction, the LOCO experiment (Section 2.5) was repeated on the corresponding batch-corrected CLR matrices, and within-cohort AUC was re-computed under the same nested-CV procedure as Section 2.4. Because Design 2’s correction step is explicitly given each sample’s own diagnosis label as part of a jointly-fit model, within-cohort AUC on Design-2-corrected features cannot be used to argue that disease-associated biological signal survived correction independent of that label information — preserved separability could instead reflect the correction algorithm’s own label-preservation mechanism rather than biology. This caveat does not apply to Design 1, whose within-cohort AUC is not label-confounded, though it remains transductive like the LOCO evaluation itself.

2.8 SHAP Feature Importance, Direction, and Cross-Cohort Overlap Analysis

To identify which genera drove predictions in each cohort and whether these features were consistent across cohorts, we computed SHapley Additive exPlanations (SHAP) values (Lundberg and Lee, 2017) using an out-of-fold (OOF) procedure to avoid the in-sample overfitting bias of computing SHAP on the same data used to fit the model.

For each labeled cohort, SHAP values were computed using the same 10-fold outer CV structure as AUC estimation (Section 2.4): in each fold, the classifier was trained on the 9 training folds, and SHAP values were computed for the held-out fold using the training fold as background. For logistic regression, `shap.LinearExplainer` was used with the training fold as background (`feature_perturbation=“interventional”`) and the held-out fold as the evaluation set, ensuring the mean SHAP value is not algebraically forced to zero. For LightGBM, `shap.TreeExplainer` was used (`check_additivity=False`) with the same OOF structure. Values were aggregated across all 10 folds; mean absolute SHAP value per genus ($|SHAP|$) was used as the importance metric for both models and is unaffected by the direction issue below.

Direction is not defined from mean signed SHAP. For a linear model under `LinearExplainer` with interventional perturbation, the SHAP value for feature j is approximately $\beta_j \cdot (x_{ij} - \bar{x}_j^{\text{background}})$, so mean signed SHAP over an evaluation set is approximately $\beta_j \cdot (\bar{x}_j^{\text{eval}} - \bar{x}_j^{\text{background}})$ — depending on both the fitted coefficient and the mean-abundance difference between evaluation and background sets, not coefficient sign alone. Within a cohort’s own OOF folds this difference is close to sampling noise, but for cross-cohort (LOCO) SHAP the background ($N-1$ training cohorts) and evaluation set (held-out cohort) come from systematically different populations by construction — the same compositional difference PERMANOVA quantifies in Section 2.6 — so mean signed LOCO SHAP is confounded with raw cross-cohort abundance differences rather than reflecting the model’s direction. Mean signed SHAP is accordingly not used anywhere in this paper as a global, cross-cohort direction or biological-association statistic.

Coefficient-based direction and stability (logistic regression only). For logistic regression, direction is instead defined from the fitted model’s own coefficient. In each of the 10 outer OOF folds per cohort, each genus’s coefficient sign was recorded. A genus’s direction in a cohort was **stable** — positive or negative — if the coefficient carried that sign in at least 8 of 10 outer folds (specified before any results were examined); an exactly-zero coefficient counted toward neither. A genus met the **descriptive coefficient-**

sign-stability screen only if it (a) appeared in the top-20 OOF mean-|SHAP| ranking in ≥ 2 cohorts, (b) had a stable positive coefficient in ≥ 1 of those cohorts, and (c) a stable negative coefficient in ≥ 1 other — a conservative definition relative to a naive sign check on mean SHAP, since it requires the direction to be reproducible across most of the 10 outer folds, not merely to differ in one pooled mean-SHAP sign. We treat this screen as descriptive, not a formal hypothesis test: the 10 outer-CV folds within a cohort share substantially overlapping training data, so fold-level coefficient signs are not independent statistical replications, and no multiplicity-adjusted p-value is assigned to any taxon. LightGBM is a tree ensemble with no single coefficient-like, globally monotonic direction, and its SHAP-genus relationships may be nonlinear; no direction or stability screen is therefore applied to LightGBM anywhere in this paper — its results are limited to OOF mean-|SHAP| importance and the top-feature-overlap/Jaccard analysis below, neither of which requires a direction assignment.

For LOCO analyses, SHAP values were computed on the held-out test cohort from each LOCO-trained model, with the three-cohort training pool as background, and are used only as a feature-*importance* ranking (mean |SHAP|); consistent with the confound above, no cross-cohort AD/CN direction is assigned to LOCO SHAP for either model.

Cross-cohort taxonomic overlap. Pairwise taxonomic overlap was quantified using Jaccard similarity on the top-N SHAP genera ($N \in \{10, 20, 50\}$; $N=20$ primary). The reported statistic is the mean of the six unique pairwise Jaccard values among the four labeled cohorts' top-N sets ($\binom{4}{2} = 6$). Two null distributions are reported. The primary, matched-statistic null mirrors the observed statistic's construction exactly: in each of 100,000 replicates, four independent random top-N gene sets were drawn (without replacement) from the shared 396-genus pool, and the same six pairwise Jaccard similarities were averaged. Empirical one-sided $p = (1 + \# \{ \text{null} \geq \text{observed} \}) / (n_{\text{perm}} + 1)$ (Davison and Hinkley, 1997).

A second, prevalence-restricted null addresses a limitation of the shared pool: per-cohort genus prevalence varies substantially (e.g., Kazakhstan reaches $\geq 20\%$ prevalence for 356/396 genera, versus 115–131 for the three MiSeq cohorts), so a genus nearly absent from a cohort could still be drawn into that cohort's comparison set under the shared-pool null. Here, each cohort's top-N ranking and random draws were both restricted to genera reaching $\geq 20\%$ prevalence within that cohort's own AD/CN-labeled samples, with 100,000 replicates per (model, N) and every cohort confirmed to have $\geq N$ eligible genera. We report both nulls without assuming significance under one implies significance under the other. Directional flips are classified using the coefficient-stability criterion above, not this overlap analysis.

2.9 AI Tool Disclosure

Initial manuscript drafts and computational scripts were prepared with minimal assistance from Claude Sonnet 5 (Anthropic). All analytical decisions, parameter choices, experimental design, statistical interpretations, and scientific conclusions were made by the author. All code and text were reviewed, validated, and substantially revised under the author's direction.

2.10 Statistical Analysis Environment

DADA2 preprocessing (Section 2.2) was run using DADA2 v1.28 in R v4.3 and was not rerun for any analysis in this paper. The initial classifier, PERMANOVA, batch-correction, and SHAP analyses (Sections 2.4–2.8, excluding the peer-review additions below) used scikit-learn v1.5, lightgbm v4, shap v0.51, pandas v2.x, and numpy v1.x (Python 3.11), and vegan v2.6, sva v3.50 (ComBat-seq), and MMUPHin

v1.15 in R v4.3. Analyses added in response to peer review — training-only feature-selection sensitivity, label-blind batch correction, SHAP coefficient-stability screen, matched-statistic Jaccard null (Sections 2.5.1, 2.7, 2.8), Zhu 2022 repeated-CV/permutation robustness (Section 2.4.1), paired within-cohort-vs-LOCO AUC-difference bootstrap (Section 2.5.2), pairwise beta-dispersion and Chinese-MiSeq PERMANOVA sensitivity (Section 2.6), and prevalence-restricted Jaccard null (Section 2.8) — were executed in Python 3.11.15 (scikit-learn 1.5.2, lightgbm 4.6.0, shap 0.51.0, matching the original package versions) and R 4.5.1 (vegan 2.7.1, sva 3.58.0, MMUPHin 1.24.0, newer than the initial R environment, reflecting routine system updates). This mixed provenance — original DADA2/R4.3 preprocessing feeding downstream analysis combining an unchanged Python environment with an updated R4.5.1 environment — is disclosed rather than described as uniform. Figures were generated with matplotlib v3.8 or v3.11 at 300 DPI.

2.11 Ethics Statement

This study is a secondary analysis of publicly available, de-identified sequencing data and involved no direct participant recruitment, intervention, or access to identifiable information. Ethical approval and informed consent for primary data collection were the responsibility of the originating studies as reported in their publications; no new participants were recruited for the present analysis.

3. Results

3.1 Cohort Characteristics and Processing

Five cohorts comprising 599 deposited samples from three countries were assembled (Figure 1A; Table 1). After DADA2 processing and chimera removal, read retention rates varied substantially across cohorts (47.9%–78.3%; Supplementary Table S2), reflecting differences in sequencing platform, library preparation, and amplicon length; the Kazakhstan cohort (single-end, NovaSeq) had notably different quality dynamics than the paired-end MiSeq cohorts. SILVA v138.1 genus-level classification and cross-cohort union produced a unified feature matrix of 396 genera. The 90 blood microbiome (B_*) samples from Zhu 2022 were excluded prior to all analyses (Section 2.4), leaving a working dataset of 509 samples across five cohorts. After construction of the global 396-genus feature schema, multiplicative zero replacement and CLR transformation were applied independently to each of the 509 samples, at $\delta=0.65$.

For supervised classifier analyses (Phases 2–3 and 5–6), the Kim/KBASE 2022 cohort was excluded because per-sample diagnosis labels are IRB-restricted (Supplementary Note S1), leaving four labeled cohorts totaling 401 samples for binary AD/CN analyses (86 Zhuang + 171 Ling + 60 Zhu fecal binary subset + 84 Kazakhstan; Zhu 2022 MCI samples excluded).

3.2 Within-Cohort Classification Performance


Within-cohort nested cross-validation yielded AUC-ROC values broadly consistent with the original publications (Figure 2; Table 2). Logistic regression achieved AUC-ROC of 0.633 [95% CI: 0.51–0.75] for Zhuang 2018, 0.865 [0.81–0.92] for Ling 2020, 0.998 [0.99–1.00] for Zhu 2022, and 0.766 [0.66–0.86] for Kazakhstan. LightGBM yielded similar values: 0.635 [0.52–0.75], 0.862 [0.80–0.91], 0.979 [0.94–1.00], and 0.694 [0.57–0.80] respectively.

The near-perfect within-cohort AUC for Zhu 2022 fecal samples (logistic regression 0.998; LightGBM 0.979; n=60 binary) is notable given the small sample size (30 per class): with 10-fold CV on 60 samples, each test fold contains ~6 samples, so fold-level AUC estimates can be unstable. PERMANOVA (Section 3.4) shows diagnosis explains $R^2=0.094$ of within-cohort compositional variance ($F=6.01$, $p=0.0001$), a statistically detectable diagnosis-associated difference in this cohort. The high AUC should nonetheless be interpreted cautiously; the LOCO experiment (Section 3.3) provides a more informative assessment of transferability.

Robustness to fold assignment and to label permutation (Section 2.4.1). Repeating the strict training-only nested-CV partition 50 times under different pre-specified seeds gave a mean pooled-OOV AUC of 0.9955 (SD 0.0027, range 0.9889–0.9989) for logistic regression and 0.9760 (SD 0.0103, range 0.9378–0.9944) for LightGBM (Supplementary Table S8; Supplementary Figure S5), closely matching the single-partition training-only values (0.9956 and 0.9700) and showing the near-perfect AUC is not an artifact of one fold assignment — a fold-assignment stability check, not independent replication, since the underlying 60 participants are identical across all 50 repetitions. A complementary label-permutation test under the same pipeline and fixed outer-fold partition (1,000 permutations) produced a null AUC distribution centered near chance for both models (logistic regression: mean 0.473, SD 0.102, 95% range [0.280, 0.673]; LightGBM: mean 0.472, SD 0.109, 95% range [0.264, 0.688]) — well below the observed values of 0.9956 and 0.9700 (Supplementary Table S9; Supplementary Figure S5). No permutation among the 1,000 reached the observed AUC for either model, giving an empirical p-value at the resolution floor for this permutation count ($p = 1/1,001 \approx 0.000999$ for both), which we do not interpret as more precise than this floor.

The three-class analysis (AD / MCI / CN) on all 90 Zhu 2022 fecal samples using logistic regression (one-vs-rest macro-AUC under nested cross-validation) yielded macro-AUC=0.761, confirming the multi-class discriminability of this cohort’s fecal compositional data.

For Zhuang 2018, within-cohort AUC values (0.633–0.635) were modest, consistent with the small sample size and weak within-cohort diagnosis effect observed in PERMANOVA ($R^2=0.017$, $p=0.066$; Section 3.4).



pcAD = preclinical AD (amyloid-PET+, cognitively normal); SE = single-end; PE = paired-end

Figure 1

Figure 1. Study design. (A) Overview of the five 16S rRNA amplicon sequencing cohorts assembled for this study, including country of origin, sample size, sequencing mode, 16S variable region, diagnosis label availability, and analytic phases in which each cohort participates. (B) Summary of the six-phase analysis pipeline, from DADA2 processing through SHAP feature-importance and logistic-regression coefficient-stability analyses.

Table 1. Cohort characteristics.

Cohort	Country	Accession	N (analyzed)	AD	CN	MCI	Sequencing	Library	Supervised Analysis
Zhuang 2018	China	PRJNA554111	86	43	43	—	MiSeq 2×300	Paired-end	Phases 2–6
Ling 2020	China	PRJNA633959	171	100	71	—	MiSeq 2×300	Paired-end	Phases 2–6
Zhu 2022†	China	PRJNA489760	90 (fecal)	30	30	30	MiSeq 2×300	Paired-end	Phases 2–6
Kaiyrlykyzy 2022	Kazakhstan	PRJNA811324	84	41	43	—	NovaSeq	Single-end	Phases 2–6
Kim 2022	South Korea	PRJEB50447	78	18*	60	—	MiSeq 2×300	Paired-end	Phase 4 only
Total			509	232	247	30			

Amyloid-PET positive preclinical AD; per-sample labels IRB-restricted and not publicly available. †Zhu 2022 deposited 180 samples (90 fecal F₋ + 90 blood B₋); only fecal samples retained for analysis.*




Figure 2

Figure 2. Within-cohort classification performance under nested cross-validation (10-fold outer / 5-fold inner CV). Bars show out-of-fold AUC-ROC; error bars represent 95% CIs from the diagnosis-stratified, participant-level bootstrap (10,000 replicates; Section 2.5.2), matching Table 2. Dashed line = chance (AUC = 0.50). LogReg = L2-regularized logistic regression; LGBM = LightGBM.

Table 2. Within-cohort nested cross-validation performance (10-fold outer, 5-fold inner CV).

Cohort	Model	N	AUC-ROC	95% CI	Sensitivity	Specificity	Originally Reported AUC
Zhuang 2018	LogReg	86	0.633	[0.51–0.75]	0.814	0.512	~0.73 (Zhuang et al., 2018)
Zhuang 2018	LGBM	86	0.635	[0.52–0.75]	0.884	0.442	—
Ling 2020	LogReg	171	0.865	[0.81–0.92]	0.780	0.859	~0.87 (Ling et al., 2021)
Ling 2020	LGBM	171	0.862	[0.80–0.91]	0.740	0.887	—
Zhu 2022	LogReg	60	0.998	[0.99–1.00]	0.967	1.000	~0.94 (Zhu et al., 2022)
Zhu 2022	LGBM	60	0.979	[0.94–1.00]	0.900	1.000	—
Kaiyrlykyzy 2022	LogReg	84	0.766	[0.66–0.86]	0.829	0.605	~0.78 (Kaiyrlykyzy et al., 2022)
Kaiyrlykyzy 2022	LGBM	84	0.694	[0.57–0.80]	0.610	0.767	—

N=60 for Zhu 2022 binary analysis (30 AD + 30 CN fecal samples; MCI and blood samples excluded). Originally reported AUC values are approximate, derived from the primary publications using the best-performing classifier reported; direct numerical comparison is limited by differences in cohort subset, model type, and validation design. All 95% CIs in this table use the diagnosis-stratified, participant-level bootstrap of Section 2.5.2 (10,000 replicates), which states which other tables/figures use this procedure versus the original 1,000-resample bootstrap.

3.3 Cross-Cohort Generalization Failure

LOCO cross-cohort AUC-ROC values revealed substantial and variable degradation across all cohorts and model types (Figure 3A; Table 3). Mean LOCO AUC across the four cohorts was 0.640 for logistic regression and 0.647 for LightGBM — mean drops of 0.176 and 0.146 respectively from within-cohort performance. One of eight model-cohort combinations produced LOCO AUC near chance: Zhuang 2018 with logistic regression (AUC=0.504 [0.38–0.63]).

The largest absolute logistic regression AUC drop was Kazakhstan (0.766 $\rightarrow$ 0.564, Δ =0.202). Zhu 2022 showed the largest LightGBM drop (0.979 $\rightarrow$ 0.758, Δ =0.221); its logistic regression drop was also substantial (0.998 $\rightarrow$ 0.813, Δ =0.185). These results are consistent with LOCO models relying on predictive patterns learned from training cohorts that do not transfer fully to the held-out cohort, where cohort-specific compositional structure may contribute to within-cohort performance. Ling 2020 also experienced substantial LOCO collapse: logistic regression 0.865 $\rightarrow$ 0.678 (Δ =0.187); LightGBM 0.862 $\rightarrow$ 0.674 (Δ =0.188).

Formal paired comparison of the within-cohort vs. LOCO AUC difference (Section 2.5.2). A paired, diagnosis-stratified, participant-level bootstrap (10,000 replicates, identical resampled participants used for both AUC values in each replicate) formally compared the within-cohort and LOCO AUC difference for each of the eight cohort-model pairs, rather than relying on non-overlap of independently computed CIs, which is not a formal test of a nonzero difference. Five of eight pairs showed a 95% bootstrap CI on the

delta excluding zero: Ling 2020 (logistic regression $\Delta=0.187$, 95% CI [0.10, 0.28]; LightGBM $\Delta=0.188$, 95% CI [0.10, 0.28]), Zhu 2022 (logistic regression $\Delta=0.184$, 95% CI [0.07, 0.31]; LightGBM $\Delta=0.221$, 95% CI [0.11, 0.35]), and Kazakhstan logistic regression ($\Delta=0.201$, 95% CI [0.04, 0.36]). The remaining three — Zhuang 2018 logistic regression ($\Delta=0.129$, 95% CI [-0.05, 0.30], proportion $\Delta\leq 0: 0.079$), Zhuang 2018 LightGBM ($\Delta=0.070$, 95% CI [-0.10, 0.24], proportion $\Delta\leq 0: 0.214$), and Kazakhstan LightGBM ($\Delta=0.104$, 95% CI [-0.05, 0.26], proportion $\Delta\leq 0: 0.092$) — showed the same-signed (within-cohort higher) point estimate but a CI including zero, not formally significant despite the consistent direction across all eight pairs. The aggregate pattern of reduced cross-cohort performance is robust, but three of eight individual comparisons do not reach formal statistical support at this sample size (full results in Supplementary Table S10).

LOCO performance was heterogeneous across cohorts: Zhu 2022 achieved the highest LOCO AUC (logistic regression 0.813), while Zhuang 2018 collapsed to near chance. Only one of eight main model-cohort combinations fell to near chance, and five of eight exceeded 0.55, so the results demonstrate substantial and variable degradation rather than uniform failure.

Sensitivity analyses confirmed generalization failure was not attributable to Kazakhstan, Zhuang 2018, or Ling 2020 individually (Supplementary Figure S3). Excluding Kazakhstan and running LOCO on the three remaining Chinese cohorts, Zhuang 2018 remained near chance (LogReg 0.512, LightGBM 0.567), showing its generalization failure is not a Kazakhstan artifact. Excluding Zhuang 2018, LOCO on the remaining three still showed AUC substantially below within-cohort baseline for Kazakhstan (LogReg 0.599, LightGBM 0.495). Excluding Ling 2020, Kazakhstan LOCO LogReg was 0.604. These results confirm the robustness of the main finding across cohort compositions.

Training-only feature-selection sensitivity (Section 2.5.1) confirmed this LOCO generalization-failure pattern is not an artifact of the global, all-cohort genus-selection step (Section 2.2). Re-deriving the retained genus set independently within each split — using only that split’s training-side samples, with the held-out cohort contributing zero information — changed LOCO AUC by at most 0.066 in either direction (largest: Kazakhstan LightGBM $0.591\rightarrow 0.525$, $\Delta=-0.066$; Kazakhstan LogReg $0.564\rightarrow 0.609$, $\Delta=+0.045$; all others ≤ 0.031). Retained genus count under strict training-only selection ranged from 150 (Kazakhstan, a comparatively high-prevalence cohort for many genera, excluded from feature selection during its own hold-out) to 373, versus 396 under the global rule, yet no cohort-model combination changed its qualitative classification. Within-cohort nested-CV AUC under training-only selection likewise moved by -0.083 to $+0.038$ (Supplementary Table S4). The global feature-selection design, while a genuine methodological limitation in principle, did not materially inflate or deflate the magnitude of the reported cross-cohort generalization failure.

Table 3. LOCO cross-cohort AUC-ROC versus within-cohort nested CV baseline.

Test Cohort	Model	Within-Cohort AUC	LOCO AUC	LOCO 95% CI	AUC Drop
Zhuang 2018	LogReg	0.633	0.504	[0.38–0.63]	0.129
Zhuang 2018	LGBM	0.635	0.565	[0.44–0.69]	0.070
Ling 2020	LogReg	0.865	0.678	[0.59–0.75]	0.187
Ling 2020	LGBM	0.862	0.674	[0.59–0.75]	0.188
Zhu 2022	LogReg	0.998	0.813	[0.69–0.93]	0.185

Test Cohort	Model	Within-Cohort AUC	LOCO AUC	LOCO 95% CI	AUC Drop
Zhu 2022	LGBM	0.979	0.758	[0.63–0.87]	0.221
Kazakhstan	LogReg	0.766	0.564	[0.44–0.69]	0.202
Kazakhstan	LGBM	0.694	0.591	[0.47–0.71]	0.103
Mean	LogReg	0.816	0.640	—	0.176
Mean	LGBM	0.793	0.647	—	0.146

Paired-bootstrap significance for the AUC drop (Section 2.5.2; Supplementary Table S10) matches the pattern above: significant for Ling 2020 and Zhu 2022 (both models) and Kazakhstan LogReg; not significant for Zhuang 2018 (both models) or Kazakhstan LGBM.

Pairwise single-cohort transfer AUC values (Figure 3B–C; Supplementary Table S1) ranged from 0.54 to 0.88 for logistic regression. The highest was Ling 2020 → Zhu 2022 (AUC=0.88), reflecting transfer between two large Chinese fecal cohorts with overlapping compositional structure; Kazakhstan → Ling 2020 also transferred reasonably (AUC=0.74). Most other pairwise AUCs clustered between 0.54 and 0.70, consistent with limited cross-cohort transferability overall.




Figure 3

Figure 3. Cross-cohort generalization failure. (A) Within-cohort (solid bars) versus LOCO cross-cohort (hatched bars) AUC-ROC for logistic regression (blue) and LightGBM (pink). Error bars on the LOCO bars = 95% CI from the diagnosis-stratified, participant-level bootstrap (10,000 replicates; Section 2.5.2), matching Table 3; within-cohort bars are shown without error bars in this panel. (B, C) Pairwise single-cohort-train / single-cohort-test AUC heatmaps for logistic regression and LightGBM respectively. Diagonal (within-cohort) values are omitted; colour scale: red = lower AUC, green = higher AUC.

3.4 PERMANOVA Variance Decomposition

PERMANOVA on Aitchison distance across all five cohorts showed cohort identity explained 19.3% of total compositional variance ($R^2=0.193$, $F=30.13$, $p=0.0001$, 9,999 permutations; Figure 4A). On Bray-Curtis dissimilarity, cohort R^2 was 0.136 ($F=19.80$, $p=0.0001$). Cohort-associated compositional variation was substantially larger than diagnosis-associated variation in this pooled four-cohort model: in a marginal two-variable model restricted to the four labeled cohorts ($n=401$), cohort explained $R^2=0.172$ ($F=28.03$,

p=0.0001) versus diagnosis $R^2=0.014$ ($F=6.82$, $p=0.0001$) — an approximately 12-fold disparity within the same model. This is an empirical statement about relative magnitude, not a decomposition of cohort into technical and biological components: “cohort” bundles sequencing platform, library preparation, recruitment site, diagnostic ascertainment, and unmeasured dietary/lifestyle factors that PERMANOVA cannot separate (Section 2.6). It contextualizes the batch-correction results in Section 3.5, but marginal R^2 values do not by themselves establish that the cohort and diagnosis axes overlap, nor explain why any batch-correction algorithm changed classifier performance in a specific direction — discussed as a separate, more limited hypothesis in Section 4.2.

The ~12-fold disparity is sensitive to cohort composition and should not be read as a fixed, generalizable figure. Restricting the same marginal specification to the three more technically comparable cohorts — Zhuang 2018, Ling 2020, and Zhu 2022, all Chinese, paired-end MiSeq V3–V4 studies (Section 2.6) — reduced the disparity from ~12-fold to ~2.6-fold: cohort $R^2=0.0588$ ($F=10.05$, $p=0.0001$) versus diagnosis $R^2=0.0229$ ($F=7.83$, $p=0.0001$) in this $n=317$ subset (Supplementary Table S11). This subset is more comparable, not identical — recruitment site, exact diagnostic procedure, and unmeasured dietary factors still differ — but removing Kazakhstan’s single-end/NovaSeq platform and Kim/KBASE’s separate country substantially changes the cohort-to-diagnosis variance ratio. The ~12-fold figure on the full four-cohort set is therefore not a stable, generalizable property of “cohort effects versus disease effects” in gut microbiome data; it is specific to the cohorts compared, driven substantially by cohorts differing in more than country of origin.

A logistic regression classifier trained to predict cohort identity from the CLR-transformed features across all 509 samples achieved macro-AUC=0.9917 (one-vs-rest, stratified 5-fold CV; per-cohort AUC: Kazakhstan 1.000, Kim 2022 0.997, Ling 2020 0.993, Zhuang 2018 0.989, Zhu 2022 0.980) — cohort membership is thus near-perfectly recoverable from the same features used for disease classification, confirming cohort-specific information (whatever its technical or biological composition) dominates the compositional feature space.

Principal coordinates analysis (PCoA) on Aitchison distance (Figure 4B; $n=509$) visually confirmed strong cohort-level clustering, with partial overlap between Chinese cohorts (Zhuang 2018, Ling 2020) but clear separation of Kazakhstan and Zhu 2022. Kim/KBASE 2022 was clearly separated from all others along PC1, consistent with its preclinical amyloid-based staging, South Korean origin, and potentially distinct dietary patterns.

Beta-dispersion testing revealed significant heterogeneity in within-group spread across cohorts ($F=13.93$, $p=0.0001$, 9,999 permutations; Figure 4C), meaning PERMANOVA $R^2=0.193$ reflects both compositional centroid shifts and differences in within-cohort spread. Kazakhstan showed the greatest within-cohort variance; the Chinese cohorts (Zhuang 2018, Ling 2020) were more tightly dispersed. Pairwise post-hoc betadisper comparisons (Section 2.6; Supplementary Table S12) show this is not uniform across all pairs: Ling 2020 differs significantly from every other cohort (all Holm-adjusted $p\leq 0.0420$) and drives most of the significant comparisons, while Zhu 2022 vs. Zhuang 2018 (Holm $p=0.439$) and Kazakhstan vs. Kim/KBASE 2022 (Holm $p=0.439$) do not differ significantly. Within the Chinese-MiSeq-only subset, dispersion remained significant overall ($F=7.946$, $p=0.0004$): Zhuang 2018 differed significantly from Ling 2020 after Holm correction (raw $p=0.0001$, Holm-adjusted $p=0.0003$), but its difference from Zhu 2022 did not survive Holm correction (raw $p=0.026$, Holm-adjusted $p=0.052$), and Ling 2020 and Zhu 2022 did not differ ($p=0.343$). This heterogeneity is relevant to interpreting PERMANOVA effect sizes and is acknowledged in Methods and Limitations.

Within-cohort PERMANOVA revealed highly variable disease-associated compositional effects: Ling 2020 ($R^2=0.053$, $F=9.53$, $p=0.0001$) and Zhu 2022 ($R^2=0.094$, $F=6.01$, $p=0.0001$) showed significant diagnosis effects, while Zhuang 2018 ($R^2=0.017$, $p=0.066$) and Kazakhstan ($R^2=0.022$, $p=0.069$) did not reach significance at $\alpha=0.05$. These effect sizes broadly rank-order with within-cohort AUC performance, triangulating the classifier results. Zhu 2022's fecal-only $R^2=0.094$ is the highest within-cohort disease effect among the four labeled cohorts, consistent with a detectable diagnosis-associated compositional difference.




Figure 4

Figure 4. PERMANOVA variance decomposition and beta-diversity structure. (A) Marginal R^2 values from a two-variable PERMANOVA model on the four labeled cohorts ($n=401$): cohort $R^2=0.172$ and diagnosis $R^2=0.014$, shown as independent bars. These Type III marginal R^2 values are estimated independently and do not sum to 1. (B) PCoA of Aitchison distance matrix (all five cohorts; $n=509$). Cohort $R^2=0.193$ ($F=30.13$, $p=0.0001$, 9,999 permutations). Ellipses = 95% normal confidence ellipses. (C) Beta-dispersion: distance from each sample to its cohort centroid ($F=13.93$, $p=0.0001$, 9,999 permutations).

3.5 Label-Blind Transductive Batch Adjustment Does Not Improve Cross-Cohort Transferability

Batch correction was applied using ComBat-seq (raw counts) and MMUPHin (relative abundances) to the four labeled cohorts with cohort as the batch variable, under the two designs of Section 2.7: a **label-blind** design (no diagnosis information used; primary analysis) and a **label-informed** design (each sample's own true diagnosis label used; exploratory sensitivity only) (Figure 5; Table 4).

Table 4. Mean LOCO AUC across four cohorts, by correction design and model.

Design	Method	LogReg Mean LOCO AUC	LogReg Mean Drop from Within-Cohort	LGBM Mean LOCO AUC	LGBM Mean Drop
—	Uncorrected	0.640	0.176	0.647	0.146
Label-blind (primary)	ComBat-seq	0.634	0.182	0.654	0.139
Label-blind (primary)	MMUPHin	0.639	0.176	0.648	0.145

Design	Method	LogReg Mean LOCO AUC	LogReg Mean Drop from Within-Cohort	LGBM Mean LOCO AUC	LGBM Mean Drop
Label-informed (exploratory)	ComBat-seq	0.493	0.322	0.555	0.238
Label-informed (exploratory)	MMUPHin	0.612	0.203	0.582	0.210

Label-blind correction (primary result): batch correction leaves cross-cohort generalization essentially unchanged. Under label-blind ComBat-seq and MMUPHin — which use no sample’s diagnosis label — mean LOCO AUC changed by no more than 0.007 in either direction relative to uncorrected, for either model (Table 4). No individual cohort collapsed: per-cohort logistic-regression LOCO AUC under label-blind ComBat-seq ranged from 0.525 (Zhuang 2018) to 0.769 (Zhu 2022), closely tracking the uncorrected values (0.504–0.813). Neither method materially changed mean LOCO AUC or recovered cross-cohort transferability under this design.

Label-informed transductive sensitivity analysis (Supplementary Figure S4): severe degradation was specific to the label-informed design itself. Under the label-informed design — which uses every sample’s own true diagnosis label during correction and therefore cannot represent a deployable or prospectively valid comparison — ComBat-seq reduced mean logistic-regression LOCO AUC from 0.640 to 0.493, with Kazakhstan collapsing to AUC=0.308 [0.20–0.42]. The same cohort under label-blind ComBat-seq scored 0.560, closely matching its uncorrected value (0.564). This side-by-side comparison indicates the severe, cohort-specific collapse was specific to the label-informed design — consistent with a correction model explicitly told each sample’s diagnosis label (`full_mod=TRUE/covariates="diagnosis"`) constructing a label-conditional transformation with no prospective analog — and should not be read as evidence that these methods generally degrade cross-cohort performance, or of an intrinsic incompatibility between ComBat-seq’s negative-binomial framework and 16S compositional data. MMUPHin’s label-informed effect was milder (LogReg mean: 0.640→0.612) but still did not improve on label-blind or uncorrected performance.

Within-cohort AUC on label-blind-corrected data was similar to uncorrected (within ± 0.05 for most cohort–method combinations; Supplementary Table S5) — a comparison free of diagnosis-label leakage during correction, since label-blind correction never saw any sample’s diagnosis label, though it remains transductive (cohort membership and raw feature values are still used) and retains whatever cohort/population confounding is otherwise present. We do not draw the same conclusion from within-cohort AUC on label-informed-corrected data, since that comparison is additionally confounded by label use: the label-informed correction’s covariate-preservation mechanism (`full_mod=TRUE/covariates="diagnosis"`) is explicitly designed to retain diagnosis-associated structure using each sample’s true label, so preserved within-cohort separability cannot be attributed to correction-independent biological signal versus the algorithm’s own label-preservation step. Only the label-blind design supports the conclusion that “within-cohort AUC preservation rules out indiscriminate signal removal” (descriptively, per Supplementary Table S5); we do not extend this conclusion to the label-informed results.

The empirical PERMANOVA finding that cohort explains approximately 12 times more compositional variance than disease status (marginal R^2 : 0.172 vs. 0.014; Section 3.4) contextualizes why a correction method targeting cohort-associated variance might struggle to leave disease-associated variance fully

intact, but does not by itself establish that the two sources overlap, nor explain the specific pattern observed here — a label-blind correction leaving performance essentially flat while a label-informed one produced large, method- and cohort-specific movements. That pattern is more directly explained by the transductive, label-informed design itself (Section 2.7) than by the general cohort/diagnosis variance disparity, and we present it as such rather than as a demonstration that linear batch correction cannot separate cohort from disease signal in general.




Figure 5

Figure 5. Effect of label-blind batch correction on LOCO cross-cohort AUC-ROC (primary analysis; no diagnosis information used during correction). (A) Logistic regression and (B) LightGBM LOCO AUC for each held-out cohort under three conditions: uncorrected (green), ComBat-seq label-blind corrected (magenta), and MMUPHin label-blind corrected (salmon). Error bars = 95% bootstrap CI (1,000 non-stratified resamples on held-out test predictions; Section 2.5). Dashed line = chance (0.50). Under this label-blind design, both methods left mean LOCO AUC essentially unchanged relative to uncorrected (within 0.007 AUC for both methods and models), neither recovering nor substantially worsening cross-cohort generalization. An exploratory label-informed transductive sensitivity comparison, in which correction used each sample’s own true diagnosis label and which produced substantially different, more severe degradation for some cohorts, is reported separately in Supplementary Figure S4 and is not treated as evidence about prospective performance (Section 2.7, Section 3.5).

3.6 SHAP Analysis Reveals Cohort-Specific Taxonomic Signatures and a Descriptive Coefficient-Stability Screen

To characterize feature-level patterns underlying generalization failure, we computed out-of-fold SHAP values for within-cohort and LOCO-trained classifiers (Figure 6). The OOF design avoids the in-sample overfitting bias seen in prior SHAP analyses of microbiome classifiers, where training-set SHAP values can reflect model memorization rather than out-of-sample importance. Feature importance (mean |SHAP|) is reported for both models; direction is reported only for logistic regression, using the fitted-coefficient-based definition of Section 2.8 (not mean signed SHAP).

Minimal overlap in top predictive taxa across cohorts, assessed against two null models. Pairwise Jaccard similarity of the top-20 SHAP genera between cohort pairs averaged 0.135 for logistic regression and 0.058 for LightGBM (Figure 6C) — the mean of the six unique pairwise Jaccard values among the four labeled cohorts. Against the primary, shared-pool null (100,000 replicates, each averaging six pairwise Jaccards among four random top-20 sets drawn from the full 396-genus pool; Section 2.8), logistic regression overlap was significantly above chance (null mean 0.027 [95% CI: 0.009–0.049]; $p < 0.0001$), as was LightGBM at this N ($p = 0.0048$). A sensitivity check across $N \in \{10, 20, 50\}$ under this

null (Supplementary Table S7) shows logistic regression significant at all three thresholds, while LightGBM was significant at N=20 and N=50 but not N=10 ($p=0.45$).

Under a more realistic, per-cohort prevalence-restricted null, the LightGBM finding does not survive. The shared-pool null draws candidate genera as if equally available in every cohort, but per-cohort prevalence varies substantially (Section 2.8). Restricting each cohort's top-N ranking and random draws to genera reaching $\geq 20\%$ prevalence within that cohort's own AD/CN samples (eligible universe: 131 Zhuang 2018, 126 Ling 2020, 115 Zhu 2022, 356 Kazakhstan; 100,000 replicates per model/N) changes the LightGBM result: overlap was **not significant at N=10, 20, or 50** (observed mean Jaccard 0.018/0.063/0.167 vs. null 0.026/0.052/0.142; $p=0.80, 0.24, 0.061$; Supplementary Table S13). Logistic regression remained significant at all three thresholds (observed 0.085/0.135/0.231 vs. null 0.026/0.052/0.142; $p=0.0017, 0.00003, 0.00001$). We treat this as a genuine reversal: LightGBM's apparent cross-cohort overlap is not distinguishable from chance once the null accounts for each cohort's own prevalence structure, while logistic regression's overlap is robust and remains the stronger of the two findings.

Cohort-specific top predictors, and a concrete illustration of why coefficient-based direction was needed. The most discriminative genus for logistic regression was *Subdoligranulum* in Zhuang 2018, while *Akkermansia* topped the remaining three cohorts (Zhu 2022 $|\text{SHAP}|=0.228$; Kazakhstan $|\text{SHAP}|=0.136$; Ling 2020 $|\text{SHAP}|=0.270$). Under the fitted coefficient, *Akkermansia* has a **stable positive coefficient in all four cohorts** (10/10 outer folds positive in every cohort; Supplementary Table S6) — higher CLR abundance consistently predicts higher AD log-odds, with no coefficient-based flip anywhere. This is a useful negative control: pooled mean SHAP for this genus is negative in Ling 2020 and Zhuang 2018 but positive in Zhu 2022 and Kazakhstan — an apparent sign difference that does **not** reflect the coefficient actually learned, but is instead a background-relative artifact (Section 2.8), demonstrating on real data why mean signed SHAP was inadequate and why the coefficient-based definition is used for every direction claim in this paper. For LightGBM, the top taxon per cohort was entirely distinct: *Lachnoclostridium* (Zhuang 2018), *Akkermansia* (Ling 2020), *Bacteroides* (Zhu 2022), *Castellaniella* (Kazakhstan) (Figure 6A; Supplementary Figure S2); no direction is assigned to these values (Section 2.8).

Coefficient-sign stability screen (logistic regression only). Using the pre-specified criterion (Section 2.8: stable positive coefficient sign in $\geq 8/10$ outer folds in ≥ 1 cohort, stable negative in $\geq 8/10$ folds in ≥ 1 other, restricted to genera in the top-20 OOF mean- $|\text{SHAP}|$ ranking in ≥ 2 cohorts), **eight genera met this descriptive stability criterion** (Figure 6B): *Agathobacter*, *Bifidobacterium*, *Coprococcus*, *Dorea*, *Lactobacillus*, *NK4A214* group, *Romboutsia*, and *Ruminococcus gnavus* group. This is deliberately a descriptive screen, not a statistically confirmed set: the 10 outer-CV folds share substantially overlapping training sets, so fold-level coefficient signs are not independent replications, and no multiplicity-adjusted p-value is assigned to any taxon (Section 2.8, Limitations). It is a smaller, more conservative set than a less rigorous alternative that classifies any genus with opposite-signed pooled mean SHAP across two cohorts as a “flip” — which, applied here, would identify twelve genera; the stability requirement removes flips not reproducible across most of a cohort's own outer folds. No comparable screen is applied to LightGBM, which has no coefficient-like, globally monotonic direction and may have nonlinear SHAP–feature relationships.

Romboutsia shows the cleanest cohort-level pattern among the eight taxa meeting the screen: a stable negative coefficient (higher relative abundance predicts lower AD log-odds under the fitted model) in all three Chinese cohorts and a stable positive coefficient in Kazakhstan (9/10 folds; Supplementary Table S6)

— the one genus among the eight whose pattern cleanly separates by country in this dataset. We do not have a specific mechanistic hypothesis for *Romboutsia*.

Lactobacillus also meets the screen but its per-cohort pattern is messier: stable-negative in Zhuang 2018 only, and stable-positive in the other three cohorts (Ling 2020, Zhu 2022, Kazakhstan; Supplementary Table S6). Diet is one plausible unmeasured confounder for genus-level differences of this kind, but this dataset contains no dietary covariate data, and we do not attribute this or any other taxon’s pattern to a specific mechanism.

Akkermansia does not meet the coefficient-stability screen at all: its fitted coefficient is stable and positive in all four cohorts (10/10 outer folds in every cohort; Supplementary Table S6), despite appearing to reverse sign under the pooled mean-SHAP statistic (above). This is the clearest illustration in this dataset of why coefficient-based, fold-stability-gated direction is necessary rather than a raw pooled-SHAP sign.

LOCO SHAP analysis: feature importance only, no cross-cohort direction assigned. When SHAP values were computed for LOCO-trained classifiers applied to held-out cohorts, *Akkermansia* was the top-importance predictor for two of four held-out cohorts — Zhuang 2018 ($|\text{SHAP}|=0.470$) and Zhu 2022 ($|\text{SHAP}|=0.372$) — while *Subdoligranulum* ($|\text{SHAP}|=0.269$) topped held-out Ling 2020 and *Christensenellaceae R-7 group* ($|\text{SHAP}|=0.368$) topped held-out Kazakhstan; no genus ranked first across all four. Consistent with Section 2.8, we do not assign an AD/CN direction to these values: the LOCO SHAP background (three training cohorts) and evaluation set (held-out cohort) are drawn from systematically different populations by construction, confounding sign with the raw cross-cohort compositional difference quantified by PERMANOVA (Section 3.4) rather than reflecting a stable, model-derived direction. The absence of a single top-importance genus across all four held-out cohorts is consistent with the classifier activating whichever genera most resemble its training-cohort signature rather than a disease marker generalizing across populations.




Figure 6

Figure 6. SHAP feature importance and coefficient-based direction analysis, logistic regression. **(A)** Top-20 SHAP genera by maximum |SHAP| across cohorts. Dot size $\propto$ mean |SHAP|; feature importance only — no direction is encoded in this panel. **(B)** Fitted-coefficient direction by cohort for the eight genera meeting a pre-specified descriptive coefficient-sign-stability screen (top-20 in ≥ 2 cohorts; a stable positive coefficient sign in $\geq 8/10$ outer folds in ≥ 1 cohort and a stable negative sign in $\geq 8/10$ folds in ≥ 1 other). Cell color = median fitted logistic-regression coefficient across the 10 outer folds; cell text = fold-count and direction. Bordered cells meet the $\geq 8/10$ stability criterion. This is a descriptive screen, not a multiplicity-corrected hypothesis test. **(C)** Pairwise Jaccard similarity of top-20 SHAP taxa between cohort pairs; mean of the six unique pairwise values = 0.135, significantly above a null matched to this statistic's construction (100,000 replicates of four random top-20 sets; $p < 0.0001$). A prevalence-restricted null sensitivity analysis (Supplementary Table S13) shows logistic regression remains significant while LightGBM does not.

4. Discussion

4.1 Systematic Quantification of Generalization Failure

This study provides a systematic, quantitative analysis of cross-cohort generalization limitations in gut microbiome-based AD classifiers. Across four labeled cohorts, two model types, two cross-cohort evaluation designs, and two batch correction methods, we consistently observe the same result: classifiers trained within one cohort do not transfer reliably to another. The mean LOCO AUC drop of 0.176 for logistic regression and 0.146 for LightGBM, with one of eight model-cohort combinations near chance, demonstrates that within-cohort performance figures — from the original publications or from this study’s own nested-CV estimates — cannot be taken as evidence of a generalizable diagnostic signal. A formal paired bootstrap test (Section 3.3) supports this drop for five of eight cohort-model pairs at conventional significance; the remaining three show the same directional pattern without reaching formal significance at this sample size, a distinction we report rather than treat all eight as equally well-supported.

This finding is consistent with the broader literature. The *Bacteroides* directionality reversal documented by Jemimah et al. (2023) across North American and Chinese cohorts, contradictory reports for *Akkermansia* (Liu et al., 2019), and highly cohort-specific patterns reported elsewhere all suggest population-level heterogeneity. A recent multi-cohort integration study (Wang et al., 2026, mSystems) similarly identified cross-geographic heterogeneity as a major obstacle; the present analysis complements that work by quantifying how severely within-cohort classifiers fail on external populations, and by showing neither correction method recovered cross-cohort transferability on average.

4.2 Batch-Adjustment Results and Limits of Mechanistic Interpretation

Two distinct findings must be kept separate here.

Empirical result: cohort identity explains substantially more compositional variance than disease status (Aitchison $R^2=0.193$ across all five cohorts; marginal $R^2=0.172$ cohort vs. 0.014 diagnosis in the four-cohort two-variable model — an approximately 12-fold disparity), and cohort membership is near-perfectly recoverable from the same CLR features (macro-AUC=0.9917). This 12-fold figure is not fixed, however: restricted to the three more technically comparable Chinese paired-end MiSeq cohorts, the same comparison falls to approximately 2.6-fold (cohort $R^2=0.059$ vs. diagnosis $R^2=0.023$; Section 3.4), and the significant beta-dispersion difference among cohorts (Section 3.4) means these R^2 values reflect both centroid separation and within-cohort spread, limiting a simple centroid-distance interpretation of cohort separation. Compositional structure was strongly associated with cohort-of-origin in the evaluated datasets, but the magnitude is sensitive to cohort composition and dispersion structure, not a single stable number. Separately: batch correction, evaluated under a label-blind design using no sample’s diagnosis label, left mean LOCO AUC essentially unchanged (within 0.007 AUC of uncorrected for both methods, both models; Section 3.5) — neither recovering cross-cohort transferability nor substantially degrading it.

Hypothesis, clearly flagged as such: the much larger cohort-associated variance component provides plausible context for *why* correction targeting cohort-associated variance might struggle to isolate and remove only that variance without also touching disease-associated variance, if the two are not orthogonal in this compositional space. PERMANOVA marginal R^2 values do not establish that overlap directly, and we did not run an experiment (e.g., a simulation with known, controlled cohort/diagnosis-axis overlap) capable of testing it. We present this as a hypothesis consistent with the observed pattern, not a demonstrated mechanism, and no longer state that batch correction cannot work in general or that this analysis proves cohort and disease axes overlap.

Under the label-informed design, in which ComBat-seq and MMUPHin were told every sample’s own true diagnosis label during a joint fit across all four labeled cohorts (Section 2.7), ComBat-seq’s correction substantially worsened LOCO AUC (mean LogReg: 0.640→0.493, one cohort collapsing to AUC=0.308). The identical comparison under a label-blind design — same cohorts, features, and downstream classifiers, only omitting the diagnosis label from correction — instead reproduces essentially flat LOCO AUC (Section 3.5, Table 4), indicating the severe degradation was specific to the label-informed design rather than evidence these methods generally degrade cross-cohort performance, or of an intrinsic incompatibility between ComBat-seq’s negative-binomial framework and 16S compositional data. One hypothesis consistent with this pattern, not established by it, is that a correction model explicitly told each sample’s diagnosis label constructs a label-conditional transformation with no prospective analog. We therefore do not attribute the label-informed collapse to a property of ComBat-seq’s modeling framework, and retain those numbers only as an explicitly-labeled exploratory sensitivity illustrating what a label-aware, non-deployable correction can do (Section 3.5).

The finding that survives this correction — and that we consider the more robust and useful result for the field — is narrower than originally stated: under label-blind transductive evaluation, two widely used batch-correction methods designed for other data types (RNA-seq counts; general microbiome meta-analysis) did not recover cross-cohort transferability for AD classification in this dataset. This demonstrates ineffectiveness at the specific task, not that these methods actively destroy signal, nor a mechanistic explanation of why such correction is difficult. The field may still need fundamentally different approaches — multi-cohort training with cohort as an explicit stratification variable, domain adaptation methods, or meta-analytic frameworks that model between-cohort heterogeneity rather than attempting to remove it — but we no longer claim to have shown mechanistically why standard correction fails, only that it did not succeed under the conditions tested here.

4.3 Implications of the Zhu 2022 Paired-Sample Design

The Zhu 2022 cohort deposited both blood microbiome (B_) and fecal (F_) samples from the same 90 participants (PRJNA489760). Blood samples were excluded because blood and fecal microbiome communities are fundamentally distinct compartments; including both would introduce a sample-type confound structurally identical to a batch effect, inflating within-cohort AUC through tissue-compartment discrimination rather than disease discrimination.

With fecal-only samples (n=60 binary), the near-perfect within-cohort AUC (logistic regression 0.998, LightGBM 0.979) is harder to attribute to a single confound. The small sample size (30 per class) is relevant — each 10-fold CV test fold contains ~6 samples, so single-fold AUC estimates can be highly variable — and Zhu 2022 showed the largest detectable diagnosis-associated compositional effect ($R^2=0.094$, $F=6.01$, $p=0.0001$) of the four labeled cohorts. The repeated-CV and label-permutation analyses (Section 2.4.1, Section 3.2) show this estimate is stable to fold assignment and clearly distinguishable from a chance-label null; it may nonetheless overstate population-level performance, being derived from only 60 participants without independent replication.

The LOCO result for Zhu 2022 (logistic regression AUC=0.813, the highest LOCO value among all cohorts) is informative: classifiers trained on the other three cohorts transfer to Zhu 2022 fecal samples better than to any other cohort, suggesting the classification pattern learned here is somewhat consistent with the Chinese and Kazakh training cohorts, unlike a spurious compartment-discrimination signal that mixed-compartment pooling would have produced — a design this study avoided by excluding blood

samples at the outset (Section 2.4). We recommend future AD microbiome studies report sample types explicitly and analyze each compartment independently before pooling.

4.4 Coefficient-Sign Stability Screen: Descriptive Cross-Cohort Direction Heterogeneity

As reported in Section 3.6, eight genera met the pre-specified coefficient-sign-stability screen ($\geq 8/10$ outer folds agreeing in sign; Section 2.8) for logistic regression: *Agathobacter*, *Bifidobacterium*, *Coprococcus*, *Dorea*, *Lactobacillus*, *NK4A214* group, *Romboutsia*, and *Ruminococcus gnavus* group. We present this descriptively, not as statistically confirmed contradictions — the 10 outer-CV folds share overlapping training data, so fold-level signs are not independent replications, and we assign no multiplicity-adjusted p-values. A less conservative alternative (opposite-signed pooled mean SHAP across two cohorts, a background-relative quantity, not the coefficient) would identify twelve genera for each model; requiring reproducibility across a cohort’s own outer folds and grounding direction in the fitted coefficient is why this paper reports eight genera for logistic regression and none for LightGBM, which has no coefficient-like global direction to test.

Romboutsia shows the cleanest cohort-level pattern: stable-negative in all three Chinese cohorts, stable-positive in Kazakhstan; we offer no mechanistic hypothesis. *Lactobacillus* also meets the screen but does not cleanly separate by country (stable-negative in Zhuang 2018 only). Diet is a plausible unmeasured confounder generally, but this dataset has no dietary covariate data, so we attribute no taxon’s pattern to a specific mechanism. *Akkermansia* illustrates why the coefficient-based, stability-gated definition matters: its pooled mean SHAP appears to reverse sign between cohorts (Section 3.6), but its fitted coefficient is positive and stable in all four cohorts (Supplementary Table S6) — the model never learned an inconsistent direction; the apparent reversal is a background-relative SHAP artifact, not the fitted model.

The eight taxa meeting this screen may reflect biological population heterogeneity, technical study effects, or both; the design cannot disentangle these sources, nor does the screen establish that any taxon’s pattern exceeds what a non-independent, overlapping-fold estimate could show by chance. If some directional disagreements arise from biological population differences rather than technical artifacts, technical batch correction would not necessarily resolve them; the present design cannot determine which source dominates, consistent with Section 4.2’s finding that label-blind correction left cross-cohort AUC essentially unchanged.

4.5 Implications for Clinical Translation

These results suggest gut microbiome-based AD classifiers are not ready for clinical deployment across heterogeneous external cohorts without prospective validation. A classifier trained on Zhuang 2018 and applied to Kazakhstan achieved AUC=0.566 for logistic regression and AUC=0.493 for LightGBM — near chance and inadequate for diagnostic deployment.

We propose three prerequisites for meaningful clinical translation: (1) **Multi-cohort harmonized training data** — classifiers should train on samples from multiple geographic and demographic populations simultaneously, with cohort or population as a stratification variable, rather than training within a single cohort and expecting transfer. (2) **Prospective cross-cohort validation as a publication standard** — any future study reporting a gut microbiome AD classifier should validate performance on at least one independent cohort from a different geographic population before claiming clinical relevance; within-cohort AUC alone, however high, is an insufficient evidential standard. (3) **Covariate-adjusted models** —

diet, geographic origin, antibiotic history, and other known microbiome-modulating covariates should be measured and accounted for through study design, stratification, covariate adjustment, or sensitivity analysis, reducing the influence of population-specific compositional baselines.

Multi-cohort harmonized collection has succeeded in other biomarker domains (e.g., UK Biobank for genomic predictors, ADNI for neuroimaging). A parallel initiative — a consortium of AD microbiome studies committed to harmonized collection protocols and prospective cross-cohort sharing — would be the most direct route to a clinically actionable microbiome diagnostic.

4.6 Limitations

Several limitations of this study should be acknowledged explicitly.

Cohort search and selection: The search strategy (Section 2.1, Supplementary Note S1) identified five publicly available cohorts — not exhaustive of all existing AD microbiome studies, and additional non-Asian cohorts would strengthen the conclusions. Diagnostic-criteria comparability (clinical AD diagnosis vs. amyloid-PET positivity) adds biological heterogeneity not fully disentangled from technical batch effects.

Mixed single-end and paired-end sequencing: Kazakhstan used single-end NovaSeq reads while all other cohorts used paired-end MiSeq 2×300 bp. Single-end reads yield lower taxonomic resolution at equivalent read length due to the absence of overlap-based error correction in DADA2, and different error models were fitted per cohort — a technical asymmetry that cannot be disentangled from biological population differences here.

Kim/KBASE 2022 excluded from supervised analyses: Per-sample amyloid-PET diagnosis labels are IRB-restricted. This cohort (78 participants; 18 amyloid-positive preclinical AD, 60 CN) contributed only to PERMANOVA variance decomposition.

China overrepresentation: Three of the four labeled cohorts are Chinese, so the cross-cohort experiment is primarily a China-internal transfer (Zhuang ↔ Ling ↔ Zhu 2022) plus transfer to/from a single Central Asian cohort. A more geographically balanced design — North American, European, African populations — would be needed to generalize the conclusions further.

Partial SRA deposit for Zhu 2022: The public archive contains 180 samples representing 90 participants, a subset of the 302 analyzed in the published study (PRJNA489760), which may not be a random sample of the full study population, potentially introducing selection bias of unknown direction.

Global genus-selection feature universe: The 396-genus feature set (Section 2.2) was selected using $\geq 20\%$ prevalence pooled across all five cohorts, so the eventual held-out cohort in every LOCO/pairwise experiment contributed prevalence information. A training-only sensitivity analysis (Section 2.5.1, Section 3.3), run under the same prespecified `random_state=42` partition used throughout, re-derived genus selection independently within every split and found LOCO AUC changed by at most 0.066 — the generalization-failure conclusion is not an artifact of this design — but global selection remains a methodological limitation in principle, and we recommend training-only feature selection as the default for future cross-cohort microbiome classifier studies. Repeated cross-validation with multiple seeds was not performed for this sensitivity analysis, so 0.066 is a single-partition estimate, not one adjusted for fold-assignment variability.

PERMANOVA cannot decompose cohort into technical and biological sources, and the cohort-versus-diagnosis disparity is sensitive to which cohorts are compared. “Cohort” here is a composite variable bundling sequencing platform, library preparation, recruitment site, diagnostic ascertainment, and unmeasured dietary/lifestyle differences that PERMANOVA and betadisper cannot separate (Section 2.6). As shown in Section 3.4, the ~12-fold disparity across all four labeled cohorts fell to ~2.6-fold when restricted to the three more technically comparable Chinese paired-end MiSeq cohorts, so this ratio depends substantially on which cohorts are compared; readers should not extrapolate the ~12-fold figure elsewhere.

Jaccard overlap significance depends on the null model chosen, and one finding reverses. As shown in Section 3.6, LightGBM’s cross-cohort top-predictor overlap was significant at N=20/50 (not N=10) under a shared-genus-pool null but not significant at any tested N under a per-cohort prevalence-restricted null; we report both rather than selecting whichever supports a stronger claim. Logistic regression’s overlap was significant under both nulls at all tested N and is the more robust finding.

Batch correction: two designs, both transductive, neither prospective. As detailed in Section 2.7, both the label-blind (primary) and label-informed (exploratory) designs estimate batch-effect parameters using the eventual held-out cohort’s own feature values and cohort membership, and neither should be extrapolated to prospective deployment. The label-informed design additionally uses every sample’s true diagnosis label and is retained only as an exploratory sensitivity; its results (including the severe Kazakhstan collapse under label-informed ComBat-seq, Section 3.5) should not be read as evidence about batch correction’s effect on biological signal in general, being substantially confounded by the design itself (Section 4.2).

Repeated cross-validation is a fold-stability check, not independent replication; three of eight within-cohort-vs-LOCO comparisons are not formally significant. The 50-repetition repeated nested-CV analysis for Zhu 2022 (Section 2.4.1, Section 3.2) uses the same 60 participants every repetition, quantifying sensitivity to fold assignment, not sampling variability across independent samples; the label-permutation test in that section formally tests whether observed AUC exceeds chance. Separately, the formal paired bootstrap test of within-cohort-vs-LOCO differences (Section 2.5.2, Section 3.3) found three of eight cohort–model pairs (Zhuang 2018, both models; Kazakhstan LightGBM) did not reach a 95% CI excluding zero, despite all eight showing the same-signed drop — we report this rather than treating the aggregate mean drop as equally well-supported for every pair.

Cross-sectional design and observational data: All cohorts used cross-sectional sampling; causal inference about microbiome–AD relationships is not supported, and all language throughout refers to associations and predictive performance, not causal effects.

SHAP and coefficient-direction interpretability limitations: As detailed in Sections 2.8 and 4.4, logistic-regression direction is defined from the fitted coefficient sign, not mean signed SHAP, and is conditional on L2 regularization (see below) and CLR-feature collinearity, not a direct estimate of biological association, significance, or causal direction. The $\geq 8/10$ -fold coefficient-sign-stability screen remains descriptive rather than a formal hypothesis test — the 10 outer-CV folds are not independent replications, and we apply no multiplicity correction or per-taxon p-value — so the eight genera meeting it should be treated as empirical patterns for targeted, covariate-adjusted, independently-replicated follow-up, not confirmed mechanistic findings. LightGBM has no coefficient-like global direction and receives no directional screen. All LOCO SHAP values (either model) remain feature-importance rankings only, with

no assigned cross-cohort direction, since the LOCO SHAP background and evaluation set come from systematically different populations, confounding sign with raw cross-cohort compositional differences.

Logistic regression models use L2 (ridge), not elastic-net, regularization: The logistic-regression models throughout (`solver="saga"`) use `penalty="l2"`, not `"elasticnet"`; an `l1_ratio` parameter is present in the estimator configuration but silently ignored by scikit-learn under the L2 penalty. Consequently, essentially no fitted coefficient is exactly zero in the coefficient-stability analysis (Section 2.8), since L2 regularization does not perform the sparse variable selection that elastic-net or L1 penalties would.

5. Conclusion

We conducted a systematic cross-cohort generalization evaluation of gut microbiome-based Alzheimer’s disease classifiers spanning five independent cohorts from three countries. Within-cohort classifiers achieved AUC-ROC of 0.633–0.998, consistent with published reports; the near-perfect Zhu 2022 result was stable across 50 repeated cross-validation partitions (mean AUC=0.9955 logistic regression, 0.9760 LightGBM) and distinguishable from a label-permutation null. LOCO AUC fell to 0.504–0.813 — a mean drop of 0.176 (logistic regression) and 0.146 (LightGBM); a formal paired bootstrap supported this drop for five of eight cohort–model pairs, with the remaining three showing the same direction without reaching formal significance at this sample size (Section 3.3). This pattern was robust to a strict training-only re-derivation of feature selection (LOCO AUC changed by at most 0.066). Batch correction under a label-blind design left mean LOCO AUC essentially unchanged (within 0.007 AUC for both methods/models) without recovering transferability; an exploratory label-informed design (which cannot estimate prospective performance) produced large, cohort-specific degradation attributable substantially to the design itself rather than the correction algorithms. PERMANOVA showed cohort identity explained substantially more compositional variance than disease status (marginal R^2 : 0.172 vs. 0.014, ~12-fold, across the four labeled cohorts) — not decomposable by PERMANOVA into technical and biological sources, and not by itself explaining the batch-correction results — falling to ~2.6-fold (R^2 =0.059 vs. 0.023) restricted to three more technically comparable cohorts, with pairwise dispersion driven predominantly by one cohort. A cohort-identity classifier achieved macro-AUC=0.9917, indicating strong cohort structure in the feature space, and generalization failure persisted across all three prespecified cohort-exclusion sensitivity configurations. Out-of-fold SHAP analysis found low taxon overlap between cohort-specific predictive signatures (mean Jaccard 0.135 logistic regression, 0.058 LightGBM); logistic regression overlap remained significant under both a matched-statistic and a prevalence-restricted null, while LightGBM was significant only under the shared-pool null — a genuine reversal, reported as such. Using a pre-specified descriptive coefficient-sign-stability screen, eight genera showed opposite, stable fitted-coefficient direction across cohorts for logistic regression; no comparable screen applies to LightGBM. These findings demonstrate substantial cross-cohort generalization limitations among the evaluated 16S datasets, and the two batch-correction methods tested under label-blind design did not recover transferability. Multi-cohort harmonized prospective studies and prospective cross-cohort validation are needed before microbiome-based AD biomarkers can achieve clinical utility.

Data Availability Statement

All raw sequencing data analyzed in this study are publicly available from NCBI SRA or ENA under accession numbers PRJNA554111 (Zhuang 2018), PRJNA633959 (Ling 2020), PRJNA489760 (Zhu 2022), PRJNA811324 (Kaiyrlykyzy 2022), and PRJEB50447 (Kim 2022). Analysis code, processed data tables, and figure generation scripts are available at <https://github.com/carhanc/microbiome-ad-generalization>.

Author Contributions

A.C. conceived the study, developed and executed the complete analysis pipeline, interpreted the results, and wrote the manuscript.

Funding

The author declares that no external funding was received for the conduct of this research.

Acknowledgements

The author thanks the original study teams of all five cohorts for depositing their data in publicly accessible archives, without which this analysis would not have been possible.

Generative-AI tool disclosure: Claude Code (Anthropic; version 2.1.104), running the Claude Sonnet 5 model (Anthropic; source: <https://www.anthropic.com>), provided minimal assistance with initial manuscript drafting and computational scripting, as detailed in Methods 2.9. All analytical decisions, parameter choices, experimental design, statistical interpretations, and scientific conclusions were made by the author.

Conflict of Interest

The author declares no conflict of interest.

References

- Callahan, B.J., McMurdie, P.J., Rosen, M.J., Han, A.W., Johnson, A.J.A., Holmes, S.P. (2016). DADA2: High-resolution sample inference from Illumina amplicon data. *Nature Methods*, 13(7), 581–583. <https://doi.org/10.1038/nmeth.3869>
- Cryan, J.F., O’Riordan, K.J., Cowan, C.S.M., Sandhu, K.V., Bastiaanssen, T.F.S., Boehme, M., et al. (2019). The microbiota–gut–brain axis. *Physiological Reviews*, 99(4), 1877–2013. <https://doi.org/10.1152/physrev.00018.2018>
- Davison, A.C., Hinkley, D.V. (1997). *Bootstrap Methods and Their Application*. Cambridge University Press. <https://doi.org/10.1017/CBO9780511802843>

Harach, T., Marungruang, N., Duthilleul, N., Cheatham, V., Mc Coy, K.D., Frisoni, G., et al. (2017). Reduction of Abeta amyloid pathology in APPPS1 transgenic mice in the absence of gut microbiota. *Scientific Reports*, 7(1), 41802. <https://doi.org/10.1038/srep41802>

Jemimah, S., Chabib, C.M.M., Hadjileontiadis, L., AlShehhi, A. (2023). Gut microbiome dysbiosis in Alzheimer's disease and mild cognitive impairment: a systematic review and meta-analysis. *PLOS ONE*, 18(5), e0285346. <https://doi.org/10.1371/journal.pone.0285346>

Jung, J.H., Kim, G., Byun, M.S., Lee, J.H., Yi, D., Park, H., et al. (2022). Gut microbiome alterations in preclinical Alzheimer's disease. *PLOS ONE*, 17(12), e0278276. <https://doi.org/10.1371/journal.pone.0278276>

Kaiyrlykyzy, A., Kozhakhmetov, S., Babenko, D., Zholdasbekova, G., Alzhanova, D., Olzhayev, F., et al. (2022). Study of gut microbiota alterations in Alzheimer's dementia patients from Kazakhstan. *Scientific Reports*, 12, 15115. <https://doi.org/10.1038/s41598-022-19393-0>

Ke, G., Meng, Q., Finley, T., Wang, T., Chen, W., Ma, W., et al. (2017). LightGBM: A highly efficient gradient boosting decision tree. In *Advances in Neural Information Processing Systems* 30, pp. 3146–3154.

Kowalski, K., Mulak, A. (2019). Brain-gut-microbiota axis in Alzheimer's disease. *Journal of Neurogastroenterology and Motility*, 25(1), 48–60. <https://doi.org/10.5056/jnm18087>

Ling, Z., Zhu, M., Liu, X., Shao, L., Cheng, Y., Yan, X., et al. (2021). Structural and functional dysbiosis of fecal microbiota in Chinese patients with Alzheimer's disease. *Frontiers in Cell and Developmental Biology*, 8, 634069. <https://doi.org/10.3389/fcell.2020.634069>

Liu, P., Wu, L., Peng, G., Han, Y., Tang, R., Ge, J., et al. (2019). Altered microbiomes distinguish Alzheimer's disease from amnesic mild cognitive impairment and health in a Chinese cohort. *Brain, Behavior, and Immunity*, 80, 633–643. <https://doi.org/10.1016/j.bbi.2019.05.008>

Lundberg, S.M., Lee, S.I. (2017). A unified approach to interpreting model predictions. In *Advances in Neural Information Processing Systems* 30, pp. 4765–4774.

Ma, S., Shungin, D., Mallick, H., Schirmer, M., Nguyen, L.H., Kolde, R., et al. (2022). Population structure discovery in meta-analyzed microbial communities and inflammatory bowel disease using MMUPHin. *Genome Biology*, 23(1), 208. <https://doi.org/10.1186/s13059-022-02753-4>

Martín-Fernández, J.A., Barceló-Vidal, C., Pawlowsky-Glahn, V. (2003). Dealing with zeros and missing values in compositional data sets using nonparametric imputation. *Mathematical Geology*, 35(3), 253–278. <https://doi.org/10.1023/A:1023866030544>

Oksanen, J., Simpson, G.L., Blanchet, F.G., Kindt, R., Legendre, P., Minchin, P.R., et al. (2024). *vegan: Community Ecology Package*. R package version 2.6-8. <https://CRAN.R-project.org/package=vegan>

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.

Plovier, H., Everard, A., Druart, C., Depommier, C., Van Hul, M., Geurts, L., et al. (2017). A purified membrane protein from *Akkermansia muciniphila* or the pasteurized bacterium improves metabolism in obese and diabetic mice. *Nature Medicine*, 23(1), 107–113. <https://doi.org/10.1038/nm.4236>

Vogt, N.M., Kerby, R.L., Dill-McFarland, K.A., Harding, S.J., Merluzzi, A.P., Johnson, S.C., et al. (2017). Gut microbiome alterations in Alzheimer's disease. *Scientific Reports*, 7(1), 13537.

1100 <https://doi.org/10.1038/s41598-017-13601-y>

1101 Wang, J., Liu, H., Lai, H., Bao, Y., Shen, M., Li, C., et al. (2026). Identifying candidate gut microbiota
1102 indicators for Alzheimer's disease through integrated data. *mSystems*, 11(4).
1103 <https://doi.org/10.1128/msystems.01754-25>

1104 World Health Organization. (2026). *Dementia*. WHO Fact Sheet. [https://www.who.int/news-room/fact-](https://www.who.int/news-room/fact-sheets/detail/dementia)
1105 [sheets/detail/dementia](https://www.who.int/news-room/fact-sheets/detail/dementia)

1106 Zhang, Y., Parmigiani, G., Johnson, W.E. (2020). ComBat-seq: Batch effect adjustment for RNA-seq count
1107 data. *NAR Genomics and Bioinformatics*, 2(3), lqaa078. <https://doi.org/10.1093/nargab/lqaa078>

1108 Zhu, Z., Ma, X., Wu, J., Xiao, Z., Wu, W., Ding, S., et al. (2022). Altered gut microbiota and its clinical
1109 relevance in mild cognitive impairment and Alzheimer's disease: Shanghai Aging Study and Shanghai
1110 Memory Study. *Nutrients*, 14(19), 3959. <https://doi.org/10.3390/nu14193959>

1111 Zhuang, Z.Q., Shen, L.L., Li, W.W., Fu, X., Zeng, F., Gui, L., et al. (2018). Gut microbiota is altered in
1112 patients with Alzheimer's disease. *Journal of Alzheimer's Disease*, 63(4), 1337–1346.
1113 <https://doi.org/10.3233/JAD-180176>
